


Neural networks

An introduction

Single-layer perceptrons


input





output


Example

Multiple layers













$$x_i^{(n)} = \sum_j w_{ij}^{(n-1)} y_j^{(n-1)}$$



$$x^{(n)} = w^{(n-1)} \cdot y^{(n-1)}$$








$$x_i^{(n)} = \sum_j w_{ij}^{(n-1)} y_j^{(n-1)}$$




$$x^{(n)} = w^{(n-1)} \cdot y^{(n-1)}$$


Activation functions








XOR network



Bias neurons

bias neuron



Do we need more than 1 hidden layer?

Universal approximation theorem

Any continuous function on a compact interval
can be approximated by a feed-forward neural
network with a single hidden layer

For any f and ε there exist a , b , and c such that

$$g(x) = \sum_i c_i \varphi(a_i \cdot x + b_i)$$

$$|g(x) - f(x)| < \varepsilon$$

for all x in the interval

Training neural networks:

Backpropagation

Initializing weights

Draw from normal distribution



$$w_{ij}^{(n)} \sim N(0, \sigma)$$

More inputs \rightarrow Less variance

$$\sigma^2 = 1/N^{(n)}$$

This prevents saturation

Gradient descent



Automatic differentiation (chain rule)


Data sets

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9

Convolutional networks and deep learning




convolution +
nonlinearity



max pooling

convolution + pooling layers



vec

fully connected layers

Nx binary classification